

X33FCON 2026

THE ART OF EVASION



Fabian Mosch

AGENDA


- ▶ Whoami
- ▶ Shellcode loader Evasion basics
- ▶ Optional or overkill features
- ▶ Machine learning based detections
- ▶ Fooling LLM static malware analysis
- ▶ Conclusion



01

WHOAMI

x33fcon 2026 – The Art of Evasion



WHOAMI



- ▶ Head of Offensive Services @r-tec
 - ▶ >10 years Offensive Security
- ▶ Founder of MSec Operations UG
 - ▶ AV/EDR Evasion for Pentesters & Red Teams
- ▶ Publishing Tools/Scripts on Github, Blogposts, YouTube-Videos
 - ▶ S3cur3Th1sSh1t / @ShitSecure




an **accompio** company



02

SHELLCODE LOADER EVASION BASICS




SHELLCODE LOADER EVASION BASICS - MINIMUM

01	Polymorphism / Metamorphism	01
02	String Encryption	02
03	Anti-Emulation	03
04	Anti-Sandbox & Cloud Analysis	04
05	Runtime Encryption / Encoding ↔ Decryption / Decoding	05

SHELLCODE LOADER EVASION BASICS

Polymorphism:




The screenshot shows a code editor window titled "polymorph_loader.nim" with a "Nim 2.0" version indicator. The code is written in Nim and implements a polymorphic shellcode loader. It features a series of conditional checks and obfuscated strings. Two specific lines are highlighted with red boxes: "discard calcHard()" and two variable declarations using the "obf" function for obfuscation. The code also includes a "Wait" procedure and a "type S" declaration.

```
1  import 1      discard calcHard()
2  2      dctx.init(ptrKey)
3  # Stag 3      discard calcHard()
4  let pM 4      dctx.decrypt(cast[ptr byte](ptrEncText), cast[ptr byte](ptrDecText), dataLen)
5  if pMe 5      discard calcHard()
6  6      dctx.clear()
7  # @@RA 7
8  8
9  # Stag 8
10 copyMe 9      var fzcvatzoetgyzfozx: string = obf("wbyjqtictakwxomlznxbpg")
11 10      var ydilmnmpgujehmcyinkstkn: string = obf("fgarvktvzigdvzquaukmj")
12 11
13 # @@RA 12      var znylrzeippjydlvkecqavutqp: string = obf("tukhembyyrjkdeweogtcjnoxq")
14 13
15 # Stag 14      var nyhdbojcqckznkwwwzwgovhsb: string = obf("obcvzcauaywpdcmbreqbyevpjh")
16 var ol 15
17 Virtua 16      var mcefosuixilpsp: string = obf("dxzxfthucfmrzdtkcojgepaqbwehioznuwzmrzoapjebpasrtvgvqyeglfcvrsribjpeco")
18 17      var uupiyjjljhemiwmeqcghxf: string = obf("vuogscigoyqsmqxczcxz")
19 # @@RA 18
20 18
21 # Stag 19      proc Wait(): void =
22 type S 20      WaitForSingleObject(-1, -1)
23 cast[S 21
24 22
```

• Nim UTF-8 • Nim UTF-8 LF x33fcon

SHELLCODE LOADER EVASION BASICS

String Encryption:



The screenshot shows a Nim code editor window titled 'encrypted_strings.nim' with the Nim 2.0 version indicator. The code implements a string encryption and obfuscation routine. A red box highlights the macro definition for 'obf', and another red box highlights the usage of 'obf' to obfuscate the 'NtQueueApcThread' string in the 'GetSyscallStub' function.

```
1  import macros, hashes
2
3  var key {.compileTime.} = hash(CompileTime & CompileDate).uint8
4
5  proc deobf*(s: string, k: uint8): string =
6      result = s
7      for i in 0..


At the bottom of the editor, the status bar shows 'Nim UTF-8 LF' and the username 'x33fcon'.


```

SHELLCODE LOADER EVASION BASICS

Anti Emulation:

- ▶ Resource Exhaustion
- ▶ Breaking the Implementation
- ▶ Lots of Alternatives¹



```
anti_emulation.nim
1  proc billionIncrements(): void =


anti_emulate.c
1  // This snippet is stolen from the reference URL below!
2  HANDLE memTes = CreateMemoryResourceNotification(0);
3  if ((int)memTes >= 40)
4  {
5      decryptAndLaunchPayload();
6  }
7  return 0;
```

¹ https://github.com/sevagas/conf/blob/main/2024/MCTTP/BypassAVDynamics_10YearsLater_extended.pdf

SHELLCODE LOADER EVASION BASICS

Anti Sandbox / Cloud:

- ▶ Disk, Memory, Processes, [...]
- ▶ Domain Join Check
- ▶ Keyboard/Mouse activity



```
1  # Credit: Idea from @Mr-Un1k0d3r
2  var MIN_COUNT = 15
3  var counter = 0
4  var current*: WCHAR
5
6  proc evade *(): void =
7      while counter <= MIN_COUNT:
8          var title: LPWSTR = newWString(256)
9          var hwnd: HWND = GetForegroundWindow()
10         if hwnd != 0:
11             GetWindowTextW(hwnd, title, 256)
12         if (title[] != current):
13             when defined(verbose):
14                 echo "Title changed to: ", $title
15                 current = title[]
16                 inc counter
17             when defined(verbose):
18                 echo "Increased counter to: ", counter
19
20
21  when isMainModule:
22      evade()
```

SHELLCODE LOADER EVASION BASICS

Encryption/Encoding:


- Only *necessary* when not decoupling the payload


A screenshot of a code editor window titled 'LoadFromDisk.nim'. The editor shows Nim code for loading a file. The code includes variable declarations for 'fileHandle', 'encString', and 'encString', a loop over 'shellcodeFile', a try-except block for file opening, and a verbose message if opening fails. It then converts the file content to a byte sequence and creates a new sequence of the same length. The editor's status bar at the bottom shows 'Nim', 'UTF-8', 'LF', and 'x33fcon' in the bottom right corner.

```
1  var fileHandle: File
2  var encString: string
3  for f in {shellcodeFile}:
4      try:
5          fileHandle = open(f, fmRead)
6          encString = fileHandle.readAll()
7          break
8      except:
9          when defined(verbose):
10             echo obf("[-] Failed to open file: ") & f
11  var encText: seq[byte] = toByteSeq(encString)
12  var decText = newSeq[byte](len(encText))
```

03

OPTIONAL OR OVERKILL FEATURES





OPTIONAL OR OVERKILL FEATURES

Publication of NimSyscallPacker:

- ▶ Includes all the minimum features
- ▶ Necessary ones plus most optional ones
- ▶ The most featureful public payload Packer/Loader
- ▶ <https://github.com/S3cur3Th1sSh1t/NimSyscallPacker/>




```
NimSyscallLoader
--> @ShitSecure
v2.2


[*] Payload length not a multiple of the BlockSize: 16
[*] Length: 282
[*] Padding payload...
[*] New Length: 288
[*] Plaintext length: 288
[*] Encrypted length: 288
[*] Hex key: 6B6D766B6F656C6762616966656368
[*] Expanded Hex key: 0x7f34e7d6fd90@[54, 66, 54, 68, 55, 54, 54, 66, 54, 70, 54, 53, 54, 67, 54, 55, 54, 50, 54, 49, 54, 57, 54, 54, 54, 53, 54, 51, 54, 56]
[*] Key length not a multiple of KeySize: 32
[*] Length: 30
```

04

MACHINE LEARNING BASED DETECTIONS



MACHINE LEARNING BASED DETECTIONS



MACHINE LEARNING BASED DETECTIONS

Scanned file

Filename

Severity

High

Actions taken

File quarantined


Objective

Tactic & technique

[Machine Learning](#) via [Sensor-based ML](#)

Specific to this detection


A file meets the machine learning-based AV protection's threshold for malicious files



MACHINE LEARNING BASED DETECTIONS




MACHINE LEARNING BASED DETECTIONS



MACHINE LEARNING BASED DETECTIONS

<https://mez0.cc/posts/evaluating-implants-with-ember/>



mez0

about posts pr

Using EMBER2024 to evaluate red team implants

26-09-2025

Introduction

A while ago I wrote a post called [Citadel 2.0: Predicting Maliciousness](#). This post was a devlog of sorts in which I discussed additions made to the project, a binary analysis framework for red-teamers: [Citadel: Binary Static Analysis Framework](#).

In the 2.0 post I introduced ML prediction capabilities to Citadel by incorporating EMBER. Originally, EMBER was the defining component in which I modelled the PE metadata extraction from. Around the same time, EMBER2024 was released which updated the model itself. As of writing this in 2025, it's the only model, paper, and codebase accessible to researchers that provides access to a complete dataset for malicious and benign samples. This is reflected in the paper where they say:

A lack of accessible data has historically restricted malware analysis research.

— EMBER2024 - A Benchmark Dataset for Holistic Evaluation of Malware Classifiers

The paper for this can be found on arxiv: [EMBER2024 - A Benchmark Dataset for Holistic Evaluation of Malware Classifiers](#)



MACHINE LEARNING BASED DETECTIONS

Ember2024 in a nutshell:

- ▶ 2,626,000 training files from VirusTotal in 2024
- ▶ Various metadata – malicious vs. non malicious
 - ▶ Including samples that were weeks later flagged as malicious
- ▶ Few MB LLM downloadable via Huggingface¹

¹ <https://huggingface.co/joyce8/EMBER2024-benchmark-models>

MACHINE LEARNING BASED DETECTIONS

EMBER2024:

- ▶ Offline testing
- ▶ Automated testing

```
c:\tools\emberanalysis>uv run ember score C:\tools\NimSyscallPacker\Loader.exe  
EMBER2024 Score: Loader.exe
```


Score	0.9990
Prediction	MALICIOUS (>= 0.9)
File Type	Win64
Model	EMBER2024_Win64.model
File Size	405,504 bytes
Entropy	7.50
Sections	11
Imports	110
Exports	0

```
c:\tools\emberanalysis>uv run ember score C:\tools\NimSyscallPacker\LoaderLowEntropy.exe  
EMBER2024 Score: LoaderLowEntropy.exe
```

Score	0.9995
Prediction	MALICIOUS (>= 0.9)
File Type	Win64
Model	EMBER2024_Win64.model
File Size	1,211,904 bytes
Entropy	5.30
Sections	11
Imports	110
Exports	0
Authenticode	No

MACHINE LEARNING BASED DETECTIONS

Ember2024 in a nutshell:



MACHINE LEARNING BASED DETECTIONS

```

Claude Code v2.1.139
Opus 4.7 (1M context) · Claude Max
/home/llm

> Take a look at /opt/tools/emberanalysis/ this script can generate
https://github.com/FutureComputing4AI/EMBER2024/tree/main and in
blog post: https://mez0.cc/posts/evaluating-implants-with-ember/
information about how it works. After any modification re-check
continue until we have a score below 0.3 at least.

```

```

• PowerShell(Set-Location "c:\tools\emberanalysis"; uv run ember sco
  Error: Exit code 2
  {
    "score": 0.9919,
    "prediction": "malicious",
    "file_type": "Win64",
    "model": "EMBER2024_Win64.model",
    "sha256": "67f71bd8e62253a33e7fe745a24d8b8300eb04f6bf1074be9e
    "basename": "baseline.exe",
    "full_path": "C:\\tools\\emberanalysis\\iat_tests\\baseline.e
    "timestamp": "2026-05-18T14:23:40.892216+00:00",
    ... +10 lines (ctrl+o to expand)

• Baseline scores 0.9919 (malicious) with 42 existing imports. Now l

```

Top 5 Score REDUCERS (best for evasion)

#	DLL	API	Score	Delta
1	user32	GetMessagePos	0.9830	-0.0064
2	kernel32	GetCommandLineW	0.9831	-0.0063
3	user32	MonitorFromWindow	0.9846	-0.0048
4	user32	ShowCursor	0.9853	-0.0041
5	user32	InSendMessage	0.9860	-0.0033

Top 5 Score INCREASERS (worst for evasion)

#	DLL	API	Score	Delta
1	kernel32	CreateProcessA	0.9985	+0.0091
2	kernel32	GetTempPathW	0.9978	+0.0084
3	kernel32	CreateThread	0.9972	+0.0078
4	kernel32	CreateProcessW	0.9971	+0.0077
5	kernel32	WriteProcessMemory	0.9968	+0.0074

b Repository:
hat was published on top for more details. Also checkout this
source code to lower the ember2024 score with all these
. Make notes for everything that worked into emberanalysis.md and

MACHINE LEARNING BASED DETECTIONS

Top 5 Score REDUCERS (best for evasion)

Final Combined Binary (38 score-reducing APIs)

Metric	Baseline	Final	Delta
Score	0.9894	0.3602	-0.6292
Prediction	MALICIOUS	BENIGN	--
Import count	42	80	+38

The combined effect is massively nonlinear – individual API deltas ranged from -0.0001 to -0.0064, but combining all 38 dropped the score by -0.6292, to benign. The model is very sensitive to the overall import table composition, not just individual entries.

#	DLL	API	Score	Delta
1	kernel32	CreateProcessA	0.9985	+0.0091
2	kernel32	GetTempPathW	0.9978	+0.0084
3	kernel32	CreateThread	0.9972	+0.0078
4	kernel32	CreateProcessW	0.9971	+0.0077
5	kernel32	WriteProcessMemory	0.9968	+0.0074

MACHINE LEARNING BASED DETECTIONS

- Massive improvement. Here are the complete results:

v2 Results: 141 Score-Reducing APIs from 12 DLLs

Version	Score	Prediction	APIs	Delta from Baseline
Baseline	0.9894	MALICIOUS	42	--
v1 Final	0.3602	BENIGN	38 added	-0.6292
v2 Final	0.0566	BENIGN	141 added	-0.9328

The score dropped from 0.9894 to 0.0566 – essentially the floor. The binary now looks more benign than most legitimate software.

New DLLs that were especially effective:

DLL	Notable Score-Reducing APIs	Best Delta
shlwapi	SHGetValueW, SHSetValueW, UrlCanonicalizeW, PathFileExistsW	-0.0119
winmm	waveOutClose, timeEndPeriod, PlaySoundW, joyGetPosEx	-0.0094

MACHINE LEARNING BASED DETECTIONS


Also relevant:

- ▶ Valid code Signing certificates
- ▶ Export Table entries
- ▶ *.pdata* section entries




05

FOOLING LLM STATIC MALWARE ANALYSIS



FOOLING LLM STATIC MALWARE ANALYSIS



21 APRIL 2026 • CYRIL FRANÇOIS • DANIEL STEPANIC • JIA YU CHAN

The Cost of Understanding: LLM-Driven Reverse Engineering vs Iterative LLM Obfuscation

Elastic Security Labs explores the ongoing arms race between LLM-driven reverse engineering and obfuscation.

🕒 23 min read 🔖 Generative AI, Detection Engineering, Malware Analysis

Introduction


Over the past few years, we have observed a significant evolution in the capabilities of LLMs to be productive and to carry out various tasks that address real-world problems, such as program synthesis, malware research, or vulnerability research. Specifically in the context of reverse engineering, LLMs are particularly effective given the right tools because they are very good at reading source code even without symbols. Not only that, thanks to their knowledge, they are capable of imitating and applying reversing methodologies.

Jump to section

- Introduction
- Key takeaways
- Claude Opus 4.6 vs Tigress
- Obfuscator benchmark
- Benchmark pipeline
- Evaluation system
- Test cases
- Results

<https://www.elastic.co/security-labs/llm-reversing-vs-llm-obfuscation>

FOOLING LLM STATIC MALWARE ANALYSIS



The screenshot shows the Claude Code application window. The top bar indicates 'Claude Code' with a tab. The main interface is split into two panes. The left pane displays a 'Welcome back' message followed by a pixelated orange robot character. Below the robot, it shows 'Opus 4.6 · Claude Max ·' and 'Organization' with a redacted name, and the current directory 'C:\tools\llmanalysis'. The right pane contains a message: 'Run /init to create a CLAUDE.md file with in...' followed by a section titled 'What's new' which lists updates: 'Added `claude agents --json` to list live Cl...', 'Added `agent\_id` and `parent\_agent\_id` attri...', and 'Status line JSON input now includes GitHub r...' with a link to '/release-notes for more'. The bottom pane shows a terminal prompt '>' followed by the user's request: 'Analyse Creator.exe using python and capstone for reversing. I want to know if this file is benign or malicious. Score it from 0 to 100 – everything above 70 would get killed in production and you are responsible for that in case something benign get's killed this could lead to disruptions. Find out what this is doing on runtime.'

```
Welcome back [REDACTED]!
```

Opus 4.6 · Claude Max · [REDACTED]
Organization
C:\tools\llmanalysis

Run /init to create a CLAUDE.md file with in...

What's new
Added `claude agents --json` to list live Cl...
Added `agent\_id` and `parent\_agent\_id` attri...
Status line JSON input now includes GitHub r...
</release-notes> for more

```
> Analyse Creator.exe using python and capstone for reversing. I want to know if this file is benign  
or malicious. Score it from 0 to 100 – everything above 70 would get killed in production and you  
are responsible for that in case something benign get's killed this could lead to disruptions.  
Find out what this is doing on runtime.
```

FOOLING LLM STATIC MALWARE ANALYSIS

- `Bash(cd "C:\tools\llmanalysis" && python << 'PYEOF'`
 `data = open('Creator.exe', 'rb').read()...)`
 `L === NIM-RELATED STRINGS ===`
 `0x024885: fatal.nim`
 `0x0248CB: system.nim`
 `... +93 lines (ctrl+o to expand)`
- Massive number of obfuscated strings found. Now let me do Capstone disassembly at the entry point and analyze the XOR/deobfuscation routines plus the dynamic API resolution.
- `Bash(cd "C:\tools\llmanalysis" && python << 'PYEOF'`
 `import pefile...)`
 `L === DISASSEMBLY AT ENTRY POINT (0x11125) ===`
 `0x11125: push    rbp`
 `0x11126: mov      rbp, rsp`

FOOLING LLM STATIC MALWARE ANALYSIS

```
• Bash(cd "C:\tools\llmanalysis" && python << 'PYEOF'  
    import pefile...)  
  L  === FUNCTION 0x16380 (String deobfuscation / hash resolver) ===  
  
    0x16380:  push    rdi  
    ... +114 lines (ctrl+o to expand)  
  
+ Flibbertigibbet... (5m 35s · ↓ 15.8k tokens)  
  L  Tip: Use /btw to ask a quick side question without interrupting Claude's current work
```

FOOLING LLM STATIC MALWARE ANALYSIS

```
import pefile...)
L === ANTI-DEBUG / EVASION CHECKS ===
    0x111BA: mov rax, qword ptr gs:[rax] ; TEB/PEB access
    0x17328: mov rax, qword ptr gs:[0x60] ; TEB/PEB access
    ... +29 lines (ctrl+o to expand)

    0x17328: mov rax, qword ptr gs:[0x60] ; TEB/PEB access
    ... +29 lines (ctrl+o to expand)

● Bash(cd "C:\tools\llmanalysis" && python << 'PYEOF'
    import pefile...)
    L === PEB ACCESS AT 0x17328 ===
        0x17300: xor     eax, 0x574bb
        0x17305: mov     r12, rcx
        ... +28 lines (ctrl+o to expand)

● Critical finding at 0x17328: this is walking the PEB->Ldr->InMemoryOrderModuleList – classic module enumeration used for stealth/evasion. Let me analyze the remaining key behaviors.
```

FOOLING LLM STATIC MALWARE ANALYSIS

- Summary


Score: 92/100 – KILL IN PRODUCTION

`Creator.exe` is a Nim-compiled malicious binary (GCC 11.1.0 / MinGW-W64, compiled today). Here's what the static analysis and Capstone disassembly behavior:


Key Malicious Behaviors Confirmed via Disassembly

1. Anti-Debug Evasion (0x343EE): Reads `PEB.BeingDebugged` via `gs:[0x60]` – alters execution if a debugger is attached
2. Stealth API Resolution (0x17328): Walks `PEB->Ldr->InMemoryOrderModuleList` to find modules by name without calling `GetProcAddress` – a signature
3. SIMD String Deobfuscation (0x16380): A 570+ XOR-operation SSE2-accelerated deobfuscation engine using 32-bit keys, called dozens of times to de runtime
4. Dynamic `Ws2_32` Loading: Explicitly loads `Ws2_32.dll` and resolves `inet_ntop` for network address operations
5. Winsock NSP Hijacking: Enumerates registry under `HKEY_LOCAL_MACHINE` and `HKEY_USERS` targeting Winsock Provider entries (`ProviderDLLs`, `InProcServ` T1547.014 persistence technique
6. Encrypted Payload: `.rdata` section has entropy of 7.95 (near theoretical max of 8.0), containing ~239KB of encrypted data that gets unpacked via runtime

Classification



FOOLING LLM STATIC MALWARE ANALYSIS



```
Claude Code v2.1.159

Welcome back [REDACTED]!

[Robot Icon]

Opus 4.8 (1M context) · Claude Max ·
[REDACTED]'s Organization
c:\tools\Optics

Tips for getting started
Run /init to create a CLAUDE.md file with instructions for Cla...

What's new
Internal infrastructure improvements (no user-facing changes)
Auto mode is now available on Bedrock, Vertex, and Foundry for...
Plugins in `.claude/skills` directories are now automatically ...
/release-notes for more

Opus 4.8 is here! Now defaults to high effort · /effort xhigh for your hardest tasks

> Check optics.exe using python and capstone for reversing. I want to know if this file is benign or malicious. Score
it from 0 to 100 - everything above 70 would get killed in production and you are responsible for that due to your
rating. In case something benign get's killed this could lead to production disruptions. █
```

FOOLING LLM STATIC MALWARE ANALYSIS

```
● The initial scan flagged some false positives in the string matching. Let me do a deeper dive on the strings and exports to understand the true nature of this binary.

● Write(deep_analysis.py)
  L Wrote 318 lines to deep_analysis.py
    1 import pefile
    2 import struct
    3 from capstone import *
    4 from collections import Counter
    5
    6 TARGET = r"c:\tools\optics2\Optics.exe"
    7
    8 def extract_all_strings(data, min_length=4):
    9     ascii_strings = []
   10     current = b""
    ... +308 lines (ctrl+o to expand)

● PowerShell(cd c:\tools\optics2; python deep_analysis.py 2>&1)
  L =====
    DEEP STRING ANALYSIS
    =====
    ... +443 lines (ctrl+o to expand)
```

FOOLING LLM STATIC MALWARE ANALYSIS

Claude Code v2.1.152

Welcome back [redacted]!

- Good – we have the full output. But I w but they're NOT in the import table) an
- Write(analyze_deep.py)
 - └ Wrote 349 lines to analyze_deep.py

```
1 #!/usr/bin/env python3
2 """
3 Deep analysis pass for Optics.ex
4 1. Dynamic API resolution detect
5 2. Cross-reference .rdata string
6 3. Compiler/framework identifica
7 4. Entry point pattern matching
8 """
9
10 import pefile
... +341 lines (ctrl+o to expand)
```

Code Pattern Analysis (Capstone)

Pattern	Count	Risk
Direct syscalls	0	Clean
int 0x2e	0	Clean
XOR (non-zeroing)	0	Clean
push+ret (ROP)	0	Clean
Indirect call (call reg)	0	Clean
Indirect jmp (jmp reg)	1	Minimal
NOP sleds	0	Clean
ROR/ROL (API hashing)	0	Clean
Embedded PEs	0	Clean
PEB walking	0	Clean

ions for Claude

olved APIs (strings reference advapi32, SHELL32, econd pass.

sted)

.

FOOLING LLM STATIC MALWARE ANALYSIS

Property	Value
SHA256	
Size	379,904 bytes (371 KB)
Arch	x86_64 (64-bit)
Subsystem	Windows CUI (Console)
Compiled	2024-05-03
Vendor	
Product	v3.9.15169.6

Runtime Behavior Assessment

Based on static analysis with Capstone disassembly and import/string/export analysis, this is what the binary does at runtime:

1. Service/Session Management – The 12 exported functions (e.g., `WorkspaceService`, `SessionManager`, etc.) indicate this is a workspace/session management service component. The code structure shows proper parameter validation, standard calling conventions, and clean function prologues.
2. Registry Reads – Queries `SOFTWARE\Workspace` (own app settings) and `SOFTWARE\Microsoft\Windows NT\CurrentVersion` (OS version). Read-only – no `RegSetValue/RegCreateKey` imports.
3. Minimal Network – Only `InternetOpenW` is imported (WININET initialization). No `connect`, `send`, `recv`, `HttpSendRequest`, or download APIs. Likely just initializes a network handle for update checks or license verification.
4. UI/Theming – Imports from `UxTheme`, `COMCTL32`, `GDI32`, `USER32` for standard Windows UI operations.
5. Process Enumeration – `Process32NextW` and `K32GetProcessMemoryInfo` suggest it monitors running processes (consistent with a workspace manager).

FOOLING LLM STATIC MALWARE ANALYSIS

Score: 25 / 100 – BENIGN

Verdict: This binary is safe for production. Do NOT kill.


This is a legitimate workspace/session management component from Ridgeline Application Group Inc. (Optics v3.9). The imports are modest and appropriate, the code is well-structured with proper security mitigations, there are zero process injection / persistence / exfiltration capabilities, and all "suspicious" string hits were false positives from harmless word-salad padding in `.rdata`. The initial automated score of 39 was inflated by these false positives; the corrected score after manual review is 25/100, well below the 70 threshold.

* Brewed for 5m 27s

* **recap:** Analyzed *Optics.exe* using Python and Capstone for reverse engineering to determine if it's malicious. Scored it 25/100 (benign) – it's a legitimate Ridgeline Application Group workspace management tool safe for production. Analysis is complete, no action needed. (disable recaps in /config)

FOOLING LLM STATIC MALWARE ANALYSIS

File Backdooring:



The screenshot shows the Claude Code interface. At the top, it displays 'Opus 4.6 · Claude Max' and 'Organization C:\tools\7z'. A status line on the right indicates a one-time notice about npm global i... and Status line commands receiving 'COLUMNS' and 'LINES' environment variables, with a link to /release-notes for more. The main area contains a user prompt: 'Check 7z.exe using python and capstone for reversing. I want to know if this file is benign or malicious. Score it from 0 to 100 – everything above 70 would get killed in production and you are responsible for that in case something benign get's killed this could lead to disruptions. Find out what this is doing on runtime.' Below the prompt, the status '★ Sublimating... (3s · thinking)' is shown.

```
Opus 4.6 · Claude Max · [redacted]  
Organization C:\tools\7z  
Claude Code now shows a one-time notice when your npm global i...  
Status line commands now receive 'COLUMNS' and 'LINES' environ...  
/release-notes for more  
Check 7z.exe using python and capstone for reversing. I want to know if this file is benign or malicious. Score  
it from 0 to 100 – everything above 70 would get killed in production and you are responsible for that in case  
something benign get's killed this could lead to disruptions. Find out what this is doing on runtime.  
★ Sublimating... (3s · thinking)
```

FOOLING LLM STATIC MALWARE ANALYSIS

File Backdooring:

- Bash(python3 << 'PYEOF'
import struct, math...)
L === SECURITY FEATURES ===
DllCharacteristics: 0x8000
ASLR (DYNAMIC_BASE): False
... +27 lines (ctrl+o to expand)
- Good data. Let me now extract the actual version info values, verify the SHA256 against known 7-Zip hashes, and check the entry point pattern more carefully.
- Bash(python3 << 'PYEOF'
import struct...)
L === VERSION INFO DETAILS ===
CompanyName : Igor Pavlov
FileDescription : 7-Zip Console
... +23 lines (ctrl+o to expand)
- Spelunking... (1m 59s · ↓ 7.5k tokens)
L Tip: Name your conversations with /rename to find them easily in /resume later

FOOLING LLM STATIC MALWARE ANALYSIS

File Backdooring:

LoadLibrary + GetProcAddress

Loads codec DLLs (7z.dll) at runtime – standard plugin architecture

Runtime Behavior Summary

This is a console-mode file archiver. At runtime it:

1. Parses command-line arguments for archive operations (compress/extract/list)
2. Reads/writes files and directories via standard Win32 file I/O
3. Uses multi-threaded LZMA/Deflate/BZip2/PPMD compression
4. Reads registry for settings/file associations (read-only)
5. Optionally loads codec DLLs for additional archive format support
6. Elevates backup/restore privileges to preserve file security descriptors
7. Outputs progress and results to the console

No network activity, no process injection, no persistence, no data exfiltration.


Maliciousness Score: 8 / 100

This is a legitimate, benign copy of 7-Zip 16.04 (console edition) by Igor Pavlov. The minor flags (missing ASLR/DEP, unsigned, .rsrc EXEC bit) are characteristics of this older open-source build, not indicators of tampering. Every imported API maps directly to expected archiver functionality. Safe for production.

* Sautéed for 3m 6s


FOOLING LLM STATIC MALWARE ANALYSIS

What do avoid:



06

CONCLUSION



CONCLUSION

- ▶ Cat & Mouse continues and probably always will
- ▶ Attackers need to get creative
 - ▶ Enhance loaders with benign metadata
 - ▶ Create benign looking payloads
- ▶ Machine learning makes evasion harder
 - ▶ But also makes evasion easier
 - ▶ All parties profit – besides from the ones not using it
- ▶ LLM reversing on the endpoint unlikely right now
 - ▶ Might however be doable in months to years

THANK YOU FOR YOUR ATTENTION!

QUESTIONS?

Fabian Mosch

accompio
we empower IT

